

# QbT: a worked example of every board page type

Specimen group, boardform board

## Abstract

This document exists to be built, not to be read for its findings. It is the manuscript half of `QbT-page-types/`, the specimen group that carries one worked page per board page type. Everything here is produced by the same tools the live papers run: the figure by its own recipe, the shipped display folder by `build-displays.py`, the bibliography by `bib-from-bank.py` from a shared bank. If any of those is broken, this document fails to compile, which is the only property it is trying to have.

## 1 What a page type owns on disk

A board page type has always been described by what a page *is*: how it resolves, how it closes, which typed records it fills. That description is complete for reading a page and incomplete for building one, because a page of some types also owns files. Ten contracts govern the types today and between them they name six artifact paths, so a writer opening a display page for the first time learns what the page must argue and not what folders it must create.

Five shapes cover the ten types. A `UNIT` type owns a folder on both sides, an authoring folder that is worked in and a shipped folder that is generated from it; display and slide are its two members. A `SECTION` type owns one flat file at the paper root and never a subdirectory of its own. A `QA-PROBE` type owns a folder of evidence records named for the page that asked; literature and value are its two members, and literature also emits a bibliography generated from a shared bank. A `MIRROR` type owns nothing and annotates a folder someone else already maintains: venue points at a shared venue playbook, design at the candidate folder its downstream display unit keeps, and skill at the skill folder it mirrors. A `NONE` type owns nothing at all, which is the honest answer for stage and meeting.

The shapes are not symmetric in who may act on them, and the display unit shows this most sharply. Figure 1 draws the eleven steps that produce one unit. Two of the eleven may be taken by a script, and both are mechanical: compiling the standalone preview, and copying the selected artifact into the shipped tree. The rest are a person's, and the three steps of the acceptance lane are a person's without exception. A machine may render a display and may ship it. No machine may accept one, because acceptance is of a specific render rather than of a name, and only a reader can say that the thing in front of them carries the claim it was asked to carry.

The practical consequence is that a contract stated in prose cannot fail, and a specimen can. A prose contract that never mentions `float.tex` reads exactly as well as one that does; a specimen folder missing `float.tex` does not build. That asymmetry is the argument for keeping a worked example of every type on disk and running the real tooling against it, rather than describing each type and trusting the description.




Figure 1: How one display unit is produced, and who is allowed to do each step. Four lanes, eleven steps, read left to right within a lane and top to bottom between them. A white box is a step a person must take; a grey box is a step a script takes. Only two of the eleven are grey, and both sit in BUILD: the standalone preview compile, and the copy into `displays/`. Every step in ACCEPT is white, which is the rule this figure exists to make visible: a machine may render a display and may ship it, and no machine may accept one. Drawn from `source/source_data.csv` at build time; the lane order is the only thing the plotting code hardcodes.

## 2 Measuring with a language model

This section exists so the literature specimen has somewhere to land. Every key below is claimed on `Qbt4-for-literature.md` and resolves in the shared bank, and the sentences say only what those sources actually say.

Large language models acquired their general few-shot behaviour at scale, which is what made them usable as measurement instruments without task-specific training (Brown et al., 2020). Applied to personality specifically, a language model has been used to augment text before a smaller model performs the detection, rather than being trusted to produce the label alone (Hu et al., 2024). Surveys of medical applications catalogue the same opportunity alongside the same reservations about reliability and evaluation (Nazi, 2024). None of this settles whether an online rating is a valid signal about a physician in the first place, which is a separate question with its own evidence: patient ratings have been compared directly against physician peer review (McGrath et al., 2018).

The point for a page-type specimen is mechanical rather than substantive. Each of the four keys was CLAIMED on a literature page with a sentence saying what job the source does, and each is cited here. The bibliography this document prints was generated from those two facts and from nothing else: no key was copied into a local file by hand, and a key that no page claims and no section cites does not appear.

## References

- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. 2020. “Language models are few-shot learners,” *Advances in neural information processing systems* (33), pp. 1877–1901.

- Hu, L., He, H., Wang, D., Zhao, Z., Shao, Y., and Nie, L. 2024. “Llm vs small model? large language model based text augmentation enhanced personality detection model,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, , vol. 38.
- McGrath, R. J., Priestley, J. L., Zhou, Y., and Culligan, P. J. 2018. “The validity of online patient ratings of physicians: analysis of physician peer reviews and patient ratings,” *Interactive Journal of Medical Research* (7:1), p. e9350.
- Nazi, K. 2024. “Large Language Models in Healthcare and Medical Domain: Applications, Opportunities, and Challenges,” *Informatics* (11:3), p. 57.  
URL <https://www.mdpi.com/2227-9709/11/3/57>